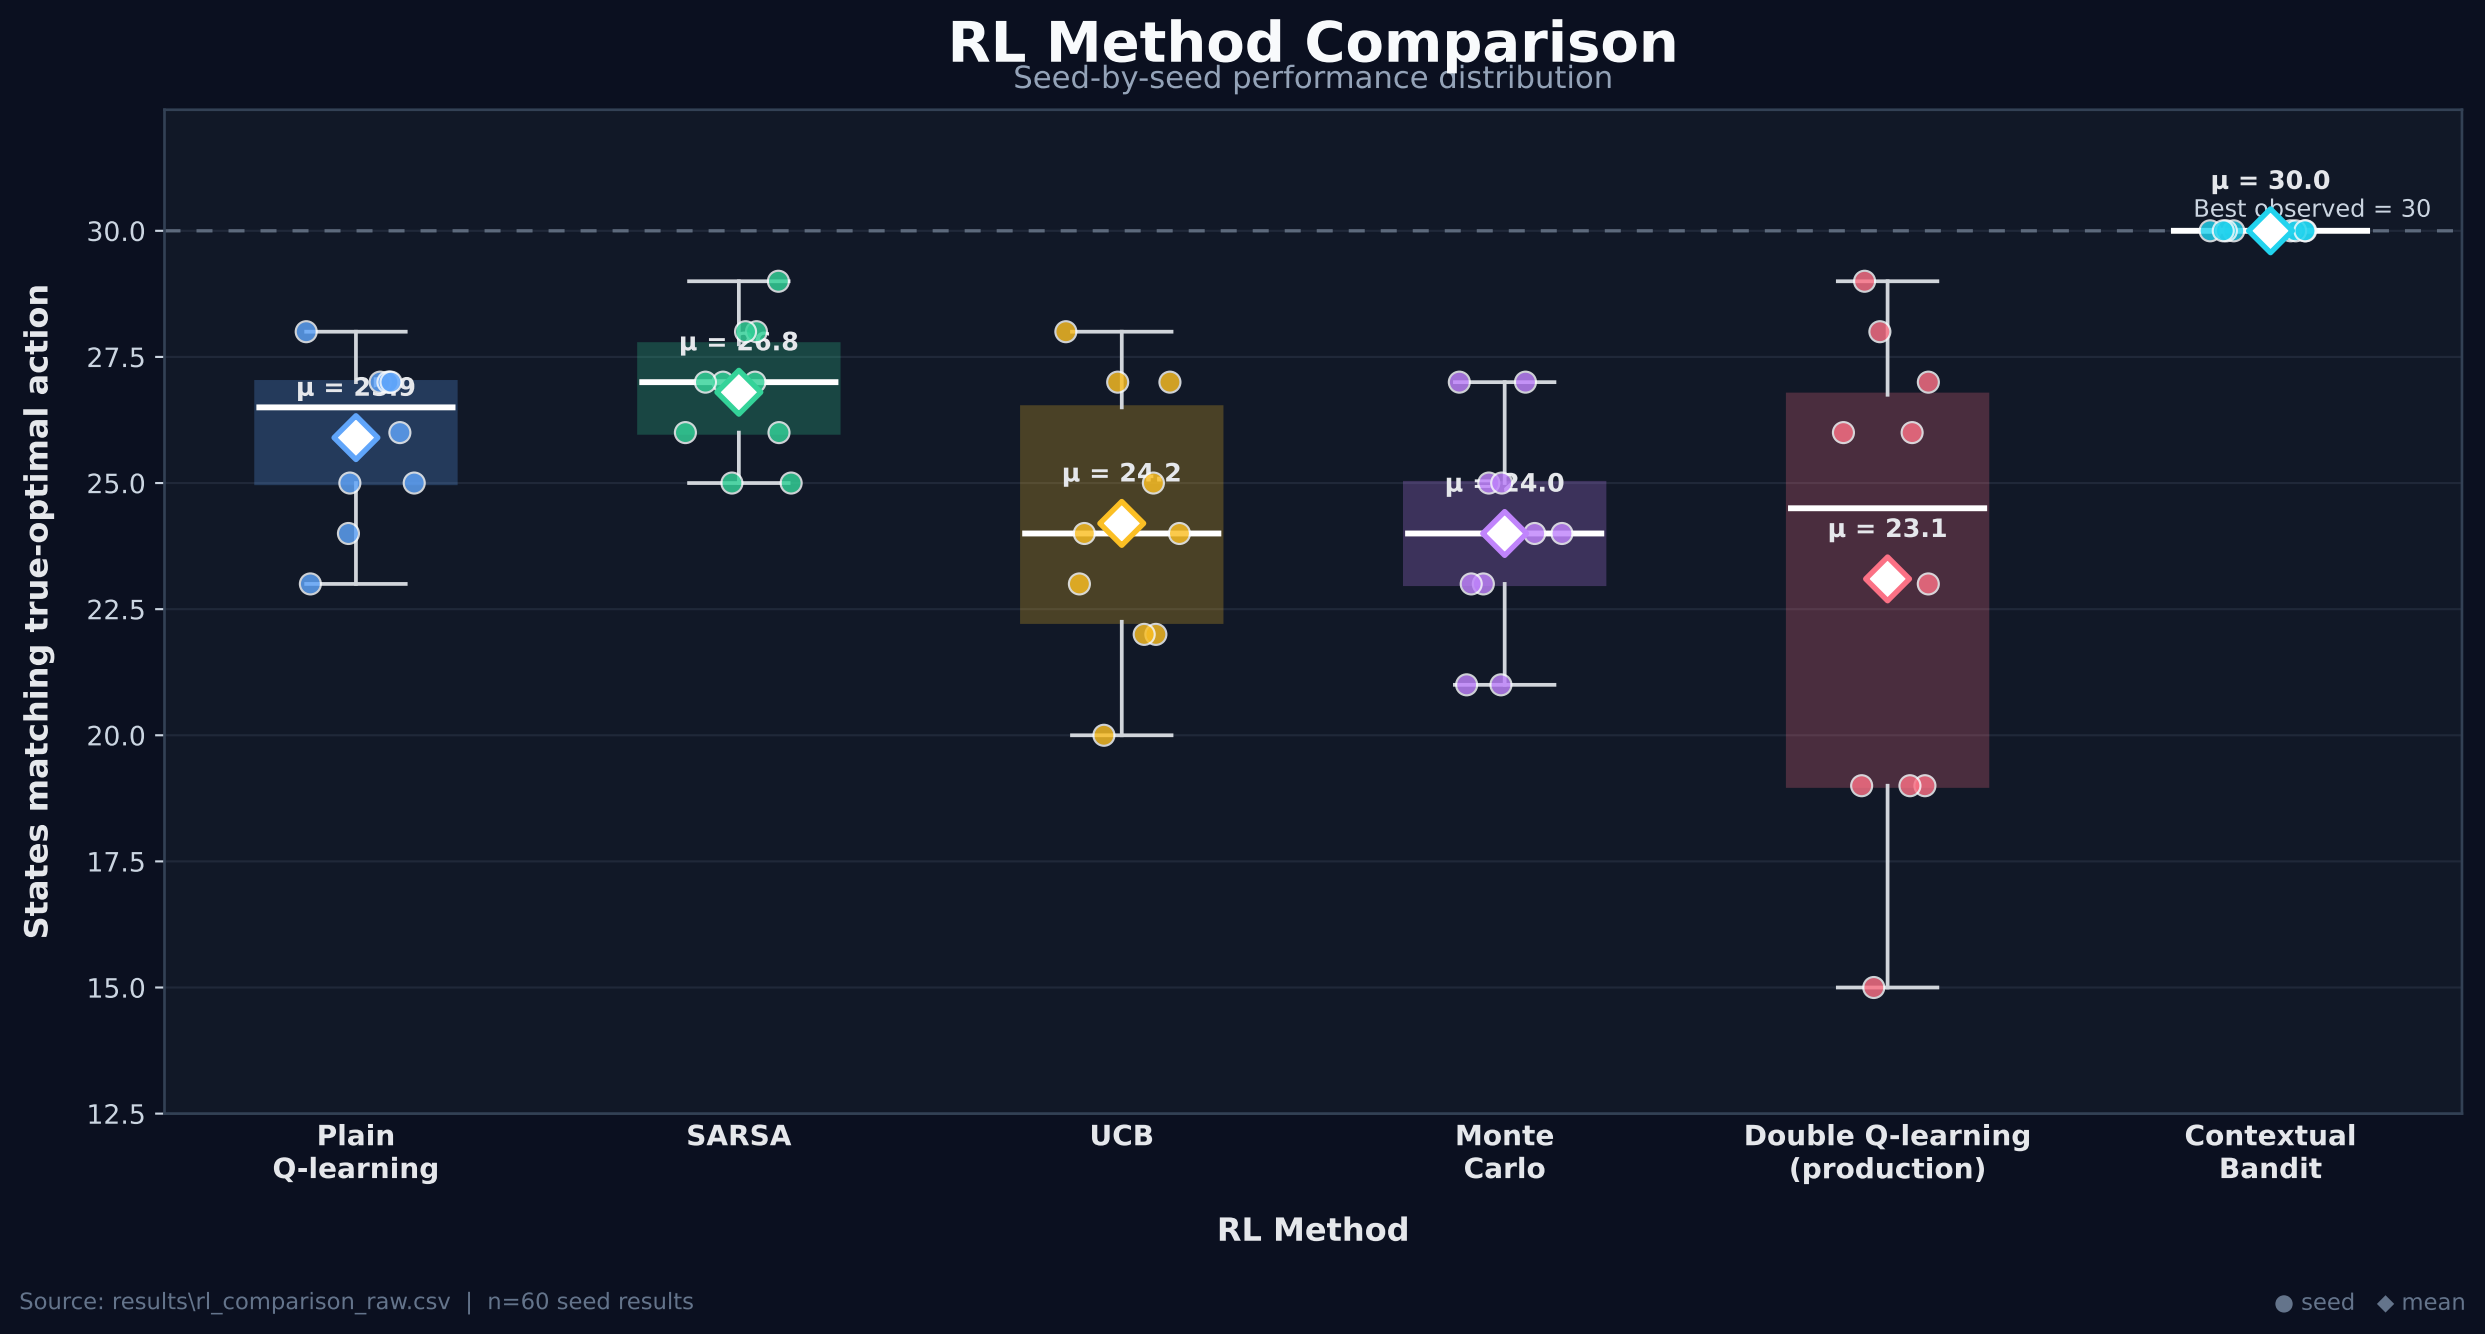
RL Method Comparison
Seed-by-seed performance distribution
States matching true-optimal action
30.0
27.5
25.0
22.5
20.0
17.5
15.0
12.5
μ = 30.0
Best observed = 30
μ = 26.8
μ = 24.2
μ = 24.0
μ = 23.1
Plain Q-learning
SARSA
UCB
Monte Carlo
Double Q-learning (production)
Contextual Bandit
RL Method
Source: results\rl_comparison_raw.csv | n=60 seed results
seed
mean